

Magic2: Supplementary Technical Reference

Architecture, Algorithms, and Implementation

Jean Thierry-Mieg

Danielle Thierry-Mieg

Greg Boratyn

April 2026 — working draft

Abstract

Magic2 is a high-throughput aligner for RNA-seq and DNA-seq reads, supporting short reads (Illumina) and long reads (Nanopore, PacBio), developed at the National Library of Medicine, NIH. This document is a technical reference describing the algorithms, data structures, parallel execution model, and output formats that constitute the Magic2 implementation. It is intended as supplementary material to the main benchmarking paper, which reports alignment accuracy and throughput comparisons against HISAT2, STAR, and minimap2; readers interested primarily in results should consult that paper. Developers and advanced users wishing to understand, extend, or port Magic2 will find the full design rationale here.

Contents

1	Overview and design philosophy	3
2	Hardware constraints and the sort-merge paradigm	3
2.1	Memory hierarchy and its consequences	3
2.2	The sort-merge alternative	3
3	Parallel execution: agents and channels	4
3.1	Channel library	4
3.2	Pipeline topology	4
3.3	NUMA-aware execution	5
4	Pilot-block inference and run-time adaptation	5
4.1	The problem: alignment parameters depend on the data	5
4.2	Pilot blocks	5
4.3	Reprocessing and zero read loss	5
4.4	Benefit	6
5	Seed generation and genome indexing	6
5.1	Seed encoding	6
5.2	Strand-agnostic seed generation	6
5.3	Paired-read encoding	6
5.4	Subsampling	7
5.5	Target multiplicity filtering and jump encoding	7
5.6	Junction seeds from annotated splice sites	7

6	Sorting, joining and optional GPU acceleration	7
6.1	Data structures	7
7	Radix sort	8
7.1	Genome index loading and joining	8
7.2	Optional GPU acceleration	8
7.3	Latency hiding through channel-based concurrency	9
7.4	Per-block GPU pipeline	9
8	Alignment extension	10
8.1	Strand resolution	10
8.2	Clustering and locus selection	10
8.3	Seed extension: the jumper algorithm	10
8.4	Junction-seed inward extension	11
8.5	Dynamic programming — best exon path	11
8.6	Adjunct-introns: splice-boundary optimisation	12
8.7	Adjunct-exons: shadow realignment	13
8.8	Alignment scoring	13
8.9	Pair scoring and multi-mapping resolution	13
8.10	Long-read rafias	13
9	Adapter detection and trimming	14
9.1	Overhang accumulation	14
9.2	Pilot blocks and reprocessing	16
9.3	Trimming and alignment extension	16
9.4	RNA-mode extras: poly(A) detection	16
9.5	Trans-splice leader detection (<i>C. elegans</i>)	16
10	Intron boundary optimisation and splice-signal scoring	16
11	Output formats	16
11.1	Alignment: SAM and BAM	16
11.2	Compact hits format	17
11.3	TSF: Tag–Sample Format	17
11.4	Coverage tracks (wiggle / BigWig)	17
11.5	Intron support tables	17
11.6	Poly(A) addition sites	18
11.7	Error profile	18
11.8	Run statistics	19
12	Automatic parameter selection	20
13	Build system and supported platforms	20
14	Current status and roadmap	20
14.1	Implemented and tested	20
14.2	Implemented, verification pending	21
14.3	Planned	21

1 Overview and design philosophy

Magic2 is a complete rewrite and extension of the earlier Magic-BLAST aligner (?), developed at the National Library of Medicine (NLM), NIH, itself building on the AceView gene and transcript annotation system (?). It aligns short (Illumina) and long (Nanopore, PacBio) RNA-seq or DNA-seq reads against an annotated or unannotated reference genome, and in a single pass produces alignments, gene-expression counts, intron-support tables, poly(A) addition sites, coverage tracks, and QC metrics.

A posteriori, the name Magic2 can be parsed as *Memory-Aware Genomic Integrative Computation* — an acronym that reflects the central role of memory efficiency in the design.

Three principles govern all architectural decisions:

1. **Cache efficiency.** The dominant cost in large-scale genomic computations is data movements, not arithmetics. All core algorithms are designed to generate long sequential memory-access patterns rather than random accesses, keeping working data inside CPU caches and saturating memory bandwidth.
2. **Automatic configuration.** Read length, paired-end geometry, adapter sequences, strandedness, and sequencing platform are inferred from the data. Internal parameters (seed length, maximum intron size, subsampling density) are derived from genome size and reported to the user for transparency. Manual overrides remain possible but are rarely necessary.
3. **Single-pass richness.** All integrative outputs — alignments, expression counts, splice-junction tables, coverage tracks, poly(A) calls, QC metrics — are produced during the alignment pass while reads and seeds reside in memory, avoiding the need for secondary tools.

The source code is available at <https://github.com/NLM-DIR/Magic2> (branch `master`). The main entry point is `wsa/sa.main.c`; the central header defining all structures and function prototypes is `wsa/sa.h`.

This document is a technical reference covering algorithms, data structures, parallel execution, and output formats. Benchmark results comparing Magic2 to current aligners are reported in the accompanying main paper; a separate user guide covers installation, index construction, and everyday usage.

2 Hardware constraints and the sort–merge paradigm

2.1 Memory hierarchy and its consequences

Modern CPUs execute each elementary operation in approximately one nanosecond when the operands reside in L1 or L2 cache, but fetching the same data from main RAM typically costs 60–100 ns on current x86 processors. A page fault to disk costs milliseconds. Because large genomic datasets far exceed cache capacity, algorithm design must prioritise access patterns over theoretical operation counts.

The core challenge in alignment is *pre-assignment*: given a short read, rapidly identify the small number of genomic loci where alignment is likely. Hash tables offer $O(1)$ average lookup, but on a multi-gigabyte genome index each lookup jumps to a quasi-random memory address, triggering cache misses and page faults. Measured latency per seed commonly reaches 100–300 ns, turning a theoretically fast algorithm into a memory-bandwidth bottleneck.

2.2 The sort–merge alternative

Magic2 adopts a *sort–merge* paradigm:

1. Generate seed records for reads and genome as large, contiguous blocks written sequentially to memory.

2. Sort both lists by seed value using a hard-coded cache-efficient radix sort (Section 7).
3. Join the two sorted lists by simultaneous linear traversal: both pointers advance monotonically, so the entire join operates strongly favors CPU cache and pre-fetching.

This approach is not novel in principle; it underlies high-performance database engines and has been applied in bioinformatics, notably in the DIAMOND protein aligner (?). Its roots lie in classical external sorting and merge-join algorithms (?), with modern refinements in cache-oblivious algorithms (?).

The key trade-off is replacing $O(n)$ cache hostile random hash look ups by two cache friendly operations: $O(n \log n)$ sorting followed by $O(n)$ sequential merging. The merge step is so cache-friendly that it runs at near-memory-bandwidth speed, far outpacing hash-table approaches on large genomes.

If available, the whole preassignment is automatically offloaded to a GPU (Section 6).

3 Parallel execution: agents and channels

3.1 Channel library

Magic2 exploits multiple cores through a clean, lock-free framework in which independent worker threads (*agents*) communicate solely through *channels* — bounded, self-synchronising queues implemented in `wchannels/channel.c`, closely following the paradigm of Go channels.

A channel is typed and has a fixed capacity. An agent inserting into a full channel blocks until space is available; an agent extracting from an empty channel blocks until data arrives. This built-in backpressure eliminates explicit locks and condition variables within agents, greatly simplifying both implementation and reasoning about correctness.

Channels track the number of active producers. When all producers have finished and decremented the source count to zero, the channel closes automatically; consumers drain any remaining records and then exit cleanly. This guarantee — no data loss and deadlock-free termination — holds regardless of the number of agents or the order in which blocks are processed.

3.2 Pipeline topology

The master thread enqueues descriptors of independent data blocks containing up to `bMax` Mega-bases (default 3 Mb) into a shared input channel. A pool of agents retrieves blocks from this channel, executes the complete per-block pipeline (seed identification, alignment extension, adaptor detection, splicing, scoring mismatches, pairing), and forwards results to designated output channels. Dedicated consumer agents drain output channels to merge the per-chromosome wiggle profiles, cumulate the gene expression counts, the introns, the poly-A sites and compile QC tables.

Each agent follows a straightforward loop:

```
while ((block = channelGet(inputChannel)) != NULL) {
    processBlock(block);
    channelPut(outputChannel, result);
}
signalDone(outputChannel);
```

Agents never wait on one another; all synchronisation emerges from channel blocking and closure. The main program defines the pipeline topology by creating channels, setting source counts, launching agents, and seeding the input channel but neither the main program nor the agents contain any explicit calls to the low level unix semaphores and other synchronization tools.

3.3 NUMA-aware execution

During its initialisation, Magic2 analyses the hardware and counts the number of CPUs, nodes and eventual GPU. Then Magic2 re-spawns itself and binds to the least-loaded node, and eventually the least busy GPU, minimising cross-node memory traffic. Users do not need set a thread count; Magic2 manages all concurrency automatically and continuously monitors its usage of memory by varying the number of datablocks processed in parallel.

For debugging, `--debug` forces fully serial execution (`nAgents=1`, `nBlocks=1`) with verbose output. `--numactl` suppresses re-spawning while retaining full parallelism.

4 Pilot-block inference and run-time adaptation

4.1 The problem: alignment parameters depend on the data

Several alignment decisions that strongly affect accuracy cannot be made before reading the data: whether the library is RNA or DNA, whether reads are short (Illumina) or long (Nanopore, PacBio), what the dominant sequencing strand is, and what adapter sequences were ligated. Requiring the user to supply these parameters manually is error-prone and unnecessary, because the data itself carries all the evidence needed to infer them.

4.2 Pilot blocks

As the input file is parsed, the read-parser emits a continuous stream of data blocks (*BB*), each carrying up to `bMax` megabases (default 3 Mb). The first five blocks are designated *pilot blocks*. They enter the agent pool and are aligned immediately, before any global parameters are frozen.

On completion of each pilot block, Magic2 tallies the evidence accumulated so far:

- **Read type.** Long reads with frequent indels indicate Nanopore or PacBio; short, low-indel reads indicate Illumina.
- **Library type.** A significant fraction of reads spanning GT-AG or CT-AC boundaries identifies the library as RNA; the absence of such signals indicates DNA.
- **Strand.** Many RNA-seq protocols are strand-specific: 99 % of reads may align consistently to one mRNA strand. The ratio of GT-AG (sense) to CT-AC (antisense) intron signals measures this directly and sets the strand bias used in all subsequent alignments.
- **Adapter sequences.** Unaligned 3' overhangs are accumulated into per-position base-frequency histograms (Section 9). Once the histogram is sufficiently populated the adapter consensus is called.

Once any single pilot block has provided enough evidence — in practice one block of 3 Mb is usually sufficient for a high-depth run — all inferred parameters are frozen and broadcast to all agents.

4.3 Reprocessing and zero read loss

If a pilot block was aligned before a parameter was finalised (for example, before the adapter was known), it is silently reprocessed from scratch with the final parameters before its results are committed. All subsequent blocks, including the remaining pilot blocks still in the channel, proceed directly with the final parameters. No reads are discarded or excluded from the output: the pilot phase is entirely transparent to the user.

4.4 Benefit

The pilot mechanism means that a single undecorated command such as

```
magic2 -i SRR122334 -x IDX -o output
```

produces fully optimised alignments regardless of library preparation, sequencing platform, or adapter chemistry. Strand bias in particular is exploited globally: a 99 % strand-specific RNA library is handled as accurately as one where the strand was declared explicitly on the command line.

5 Seed generation and genome indexing

5.1 Seed encoding

Magic2 uses exact k -mer seeds of length $k \leq 19$, encoded as 32-bit unsigned integers (2 bits per base, so 16 bases fill exactly 32 bits). The default k is chosen automatically from genome size:

- $k = 16$ for small genomes (< 1 Gb; e.g., *C. elegans*, *D. melanogaster*);
- $k = 18$ for large genomes (human, mouse).

When $k > 16$, the extra bases are handled by radix partitioning into 16 or 32 separate seed tables indexed by the least significant base pair(s), providing effective 18- or 19-mer specificity while keeping each individual table within 32 bits.

Continuous (ungapped) k -mers are preferred over spaced seeds: spaced seeds are counterproductive for indel-prone technologies (Nanopore, PacBio) because indels in the spacers disrupt the spaced-seeds.

5.2 Strand-agnostic seed generation

To halve index size and enable strand-agnostic matching, both the forward k -mer S and its reverse complement C are computed simultaneously using an incremental rolling update:

$$S_{\text{new}} = ((S \ll 2) | b) \& \text{mask} \quad (1)$$

$$C_{\text{new}} = ((C \gg 2) | (b' \ll (k - 2))) \& \text{mask} \quad (2)$$

where b is the next base (2-bit coded) and b' its complement. The strand-agnostic value $Z = \min(S, C)$ is retained. The coordinate is encoded as $(x \ll 1)$ if S was chosen (forward strand) or $(x \ll 1 | 1)$ if C was chosen (reverse strand). Ambiguous bases reset the rolling accumulator; a seed is emitted only after k consecutive unambiguous bases.

5.3 Paired-read encoding

Paired-end reads arrive with names of the form $xx/1$ and $xx/2$. Only the base name xx is entered into the read dictionary (`dictAdd(dict, "xx", &n)`), yielding a compact integer n . Then mate 1 and mate 2 are stored as contiguous numbers $2n$ and $(2n + 1)$ such that sorting by reads automatically implies sorting by pairs. The alignment stage therefore receives tables in which read-pair processing requires no further regrouping.

5.4 Subsampling

To reduce index size without sacrificing alignment sensitivity, seeds are sub-sampled using a *rolling-min-hash*: a round-robin buffer of size p (default 4) stores the hash of Z at recent positions; only the position with the minimum hash among the last p values is retained as a seed. This guarantees:

- average inter-seed spacing of $p/2$ positions;
- maximum gap of p positions;
- near-perfect synchronisation between read and genome seeds (in practice alignment begins by the second seed even when $p \leq 8$).

5.5 Target multiplicity filtering and jump encoding

After collection, target seeds are sorted (Section 7) and their multiplicity (number of occurrences in per target) is counted. Seeds with multiplicity $> Q$ (default 81) are discarded to limit spurious hits from repetitive elements; empirical tests showed $Q = 31$ gives higher specificity while $Q = 81$ slightly reduces false negatives at acceptable cost in index size and runtime.

Jump information is then embedded directly in the seed records: every 256th seed stores pointers to future seeds at distances $n_1 \times 256$, $n_2 \times 256$, $n_4 \times 256$, enabling fast skip-ahead during the merge step.

The filtered, jump-augmented target table is written to disk in compact binary format for later memory-mapped access during production alignment. Seed records are 16 bytes each (Section 6.1). overhead.

5.6 Junction seeds from annotated splice sites

When a GFF/GTF annotation is supplied, Magic2 supplements genomic seeds with *junction seeds* spanning annotated exon–exon boundaries, improving detection of short exons with known junctions.

For each annotated mRNA, all exons are concatenated *in silico* to form a continuous spliced transcript. All k -mers straddling each junction with at least 4 bases on each side ($k - 4$ starting positions per junction, with $p = 1$ so every valid position is used) are generated. For each such k -mer, two seed records with identical value Z are emitted:

- one pointing to the donor site (end of upstream exon);
- one pointing to the acceptor site (start of downstream exon), augmented with ancillary phase-shift information encoding the exact junction offset.

This dual-record strategy ensures that a single seed match near a junction can trigger extension on both sides of the splice site.

The additional index size is negligible: with $k = 18$ and $\sim 300,000$ annotated human introns, $(k - 4) \times 300,000$ yields only a few million additional seeds, compared to ~ 1 billion genomic seeds. Junction seeds are merged with the genomic table, sorted, and stored in the same binary format.

6 Sorting, joining and optional GPU acceleration

6.1 Data structures

Seed records are CW structs (16 bytes, 16-byte aligned), declared in `wsa/sa.common.h`:

Field	Type	Meaning
seed	unsigned int	32-bit 16-mer seed
nam	unsigned int	read or chrom identifier
pos	unsigned int	coordinate of first base of seed
flags	unsigned int	auxiliary information

The 16-byte size is deliberate: four CW records fill one 64-byte cache line exactly, which is optimal for both the CPU radix sort and GPU coalesced memory access.

7 Radix sort

Radix sort runs in $O(n)$ time and generates purely sequential memory accesses, making it strictly preferable to comparison-based sorts once the table exceeds cache capacity.

Magic2’s implementation (`wsa/sa.sort.c`) sorts the 32-bit `seed` field of 16-byte CW records in three passes with bit-widths of 11, 11, and 10. All three frequency histograms are accumulated in a single linear scan before any data movement, so the table is read sequentially exactly four times in total. The scatter passes ping-pong between the source array and a 128-byte-aligned scratch buffer; the sorted result lands in the scratch buffer with no final copy. On SSE2-capable processors each 16-byte record is moved by a single `_mm_load/store_si128` instruction, with a portable `memcpy` fallback elsewhere.

This design replaced an earlier merge-sort and reduced elapsed and CPU time by approximately 15 % on large human datasets. When a GPU is available the sort is offloaded to `thrust::sort` (Section 6), superseding the CPU path entirely for the seed-matching phase.

7.1 Genome index loading and joining

The filtered genome index (Section 5) is partitioned into NN sub-tables (16 or 32, indexed by the least-significant base pair(s) of the seed value). At the start of a run, all partitions are memory mapped. The read seeds and genome seed records are identical.

Both tables are scanned in parallel, advancing the pointer on the side of the lower seed. When there is a coincidence, the local Cartesian product of all entries with the matching seed is emitted as 64 bytes SEEDMATCH structs, by simply joining the 2 matching records. This table is then sorted by read (hence automatically by pairs).

7.2 Optional GPU acceleration

When an NVIDIA GPU is available, Magic2 offloads three steps of the seed-matching pipeline to the device via CUDA Thrust (`wsa/sa.gpusort.cu`): sorting the read-seed table, and joining it against the pre-loaded genome index, then sorting the table of SEEDMATCH.

GPU use is activated at compile time (`-DUSEGPU`) and detected at run time; execution falls back silently to the CPU radix sort (Section 7) when no compatible device is present. If a GPU is present, the genome seed-index is transferred to the GPU prior to the calculation of the first data block and discarded from the CPU side.

All structure declarations are shared and since both the CPU and the GPU use 64-bytes aligned buffer. no data conversion is required at the CPU–GPU boundary.

7.3 Latency hiding through channel-based concurrency

The GPU call in each agent is synchronous: the agent blocks on the device until the sorted `SEEDMATCH` table is ready, then resumes. No asynchronous CUDA streams, double-buffering, or explicit staging logic are required.

The standard GPU programming approach to this problem uses CUDA streams with asynchronous transfers (`cudaMemcpyAsync`): while the device processes block N , the host prepares and begins transferring block $N + 1$ into a second buffer, so that the PCIe transfer latency of one block is overlapped with the GPU computation of the next. This *double-buffering* pattern is effective but requires the programmer to manage two buffer sets, carefully sequence stream operations, and reason explicitly about which synchronisation points are safe — a non-trivial engineering burden that is easy to get wrong silently.

Magic2 avoids this complexity entirely. The agent pool, coordinated by bounded blocking channels, provides CPU–GPU overlap as a natural consequence of its concurrency model. While one agent is blocked waiting for the GPU, all other agents continue executing their own blocks on CPU — extracting seeds, extending alignments, writing coverage tracks, and accumulating junction and QC tables. The blocked agent simply parks at what is, from the pipeline’s perspective, an ordinary blocking-channel operation, indistinguishable from blocking on an empty input queue. Provided the number of active agents is sufficient to keep both the GPU and the CPU cores occupied — a condition that holds automatically given Magic2’s block-size and agent-count defaults — GPU transfer and sort latency is fully overlapped with productive CPU work on other blocks, with no GPU-specific orchestration code in the pipeline at all.

The channel library (`wchannels/channel.c`) is a direct C implementation of Go’s channel semantics (?), themselves rooted in Hoare’s Communicating Sequential Processes (?). It was written to bring Go’s concurrency model to an existing 700,000-line C codebase that could not be migrated to Go. Go programmers obtain the same GPU-latency-hiding property for free from the goroutine scheduler, which parks a goroutine blocked on a channel and runs another in its place. Magic2’s agent/channel architecture replicates this behaviour in C, without any GPU-specific orchestration code in the pipeline.

7.4 Per-block GPU pipeline

GPU acceleration is applied at the block level: each agent, upon receiving a block, offloads the seed-matching phase to the device, then resumes CPU-side alignment extension once the device returns. The genome seed index is resident in device memory throughout the run, having been transferred once before the first block is processed; the CPU no longer needs to memory-map the full 22 GB index during alignment.

Within each block the GPU performs the following steps in sequence:

1. **Read-seed sort.** The read seeds extracted from the current block are sorted on the device by seed value, producing the same ordered table that the CPU radix sort would deliver (Section 7).
2. **Common-seed identification.** The sorted read-seed table is intersected with the resident genome index to find all seed values present in both. For each common value the set of matching genome positions and the set of matching read positions are identified.
3. **Match-record emission.** For each common seed with m genomic occurrences and n read occurrences, all $m \times n$ candidate hit pairs are written to a match table in parallel, using a prefix-scan to assign output positions without atomic operations.
4. **Sort by read pair.** The match table is sorted by read identifier. Because of the paired-end encoding (Section 5.3), this simultaneously groups hits by read pair and places mate 1 immediately before mate 2, delivering a table directly consumable by the CPU alignment stage with no further reordering.

5. **Return to CPU.** The sorted match table is copied back to host memory. All subsequent processing — alignment extension, splicing, scoring — proceeds on the CPU without any further access to the genome index.

8 Alignment extension

The alignment extension stage receives `SEEDMATCH` records — pairs of (genome coordinate, read coordinate) for matching k -mers ($k \approx 18$), sorted by read pair. Two flavours exist: standard genomic seeds, and junction seeds which each produce two glue points (a donor and an acceptor on either side of an annotated intron boundary).

8.1 Strand resolution

Seeds are strand-agnostic: chromosome and read identifiers encode strand implicitly — odd identifiers ($2n+1$) indicate the complement was selected, even ($2n$) the forward strand. The XOR (`chrom ^ read`) & 1 therefore determines whether read and chromosome are antiparallel or parallel. This flag is extracted and dropped; from this point all processing uses the positive-strand read (`dnaR`) against either `dnaG` or `dnaGR`, both memory-mapped at initialisation.

8.2 Clustering and locus selection

For each read, glue points may scatter across many chromosomes but cluster into one or a few genuine loci. Loci are scored by glue-point count, with seeds of high genomic multiplicity (recorded in the `SEEDMATCH` record) down-weighted. Junction seeds are tallied separately and contribute additional evidence for a specific splice structure. Within each candidate locus, glue points are sorted by chromosome, strand, then by the diagonal value $a - x$ (genome coordinate minus read coordinate); in the absence of indels all seeds from the same exon share the same diagonal.

8.3 Seed extension: the jumper algorithm

Starting from each seed match, Magic2 extends the alignment using the *jumper* algorithm, first described in the AceView human-genome aligner (?) and subsequently used in Magic-BLAST (?), to which we refer for a full description.

Extension proceeds base by base: if `*cp & *cq != 0` (a bitwise test that handles IUPAC ambiguity codes), both the read pointer `cp` and the genome pointer `cq` advance by 1. On a mismatch, a *jumper table* is consulted. Each row encodes one error hypothesis as a triple (dx, da, n) : advance the read pointer by dx and the genome pointer by da , then require n consecutive matching bases for the hypothesis to be accepted. The table covers substitutions, insertions, and deletions of up to `MAXJUMP=3` bases (Table 1); a sentinel entry $\{1, 1, 0\}$ with $n = 0$ is always accepted and terminates the search without a special exit condition. Extension halts when the error density over a short window exceeds the budget defined by `errCost=4` and `errRateMax=12%`, or when the sequence ends.

The jumper algorithm is strictly linear and cache-friendly, and can drift freely along any diagonal — unlike banded Smith–Waterman, which is constrained to a fixed band and cannot follow a read that accumulates locally clustered errors. Smith–Waterman is optimal only under the assumption that sequencing errors are independently distributed, a hypothesis that does not hold in practice (polymerase slippage, homopolymer runs, and cycle-dependent quality decay are all correlated). The jumper table can be tuned to match observed error profiles; `MAXJUMP=3` is appropriate for Illumina reads and may be relaxed for Nanopore data.

Table 1: Jumper table (representative rows; fourth column not used in Magic2).

dx	da	meaning
1	1	substitution
1	0	insertion in read (1 base)
0	1	deletion in read (1 base)
2	0	insertion (2 bases)
0	2	deletion (2 bases)
$\vdots$	$\vdots$	$\vdots$
10	0	insertion (10 bases)
0	10	deletion (10 bases)
1	1	sentinel (always accepted)

8.4 Junction-seed inward extension

Junction-seed glue points are extended inward toward the exon body: the donor glue point (at the 3' end of the upstream exon) is extended $3' \rightarrow 5'$, and the acceptor glue point (at the 5' end of the downstream exon) is extended $5' \rightarrow 3'$. Together they define a candidate alignment split across the intron without ever examining the intronic sequence, which is absent from the read. A seed hit is considered redundant if it falls within ± 3 diagonals of an already-extended fragment; hits further off-diagonal are extended independently to accommodate small indels.

8.5 Dynamic programming — best exon path

Input. The input to the DP is a list of candidate exon fragments for a single read, each carrying: read and chromosome identifiers, genomic and read-space coordinates, and the error table produced by the jumper algorithm (Section 8.3). Fragments are sorted by their 5' genomic coordinate before the DP begins.

DAG construction and path extension. The DP maintains a list of active paths, initially empty. It scans candidate exons from 5' to 3' and at each step attempts to append the current exon to every existing path whose topology is compatible: the new exon must lie downstream on the same chromosome and strand, with a genomic gap consistent with an intron or with direct adjacency. When a compatible path is found, the exon is appended and the path score is updated as the running sum of individual exon scores, ignoring overlaps in read coordinates for now (those are resolved later by the adjunct-introns step, Section 8.6). A new single-exon path is also started at each exon, so that paths beginning at any position are considered.

Pruning — optimality in a DAG. The candidate set naturally forms a directed acyclic graph (DAG): multiple paths can converge on the same exon from different upstream histories. When two paths can both be extended by the same next exon, they represent two routes between a common start and a common continuation point. The path with the lower running score is pruned at this point. This pruning is optimal: because all future extensions apply identically to both paths, the lower-scoring path can never overtake the higher-scoring one regardless of what follows. Pruning keeps the active path list small and prevents combinatorial explosion on reads that cross regions with many overlapping fragment candidates.

Output and downstream refinement. At the end of the scan, the surviving paths are passed to the adjunct-introns and adjunct-exons steps (Sections 8.6–8.8), which resolve overlaps, optimise splice boundaries, realign against the genomic shadow, and compute final scores. Only then are paths ranked and the best retained. Some human genes have exact paralogues or occur in tandem repeats in the genome; reads from these loci

will genuinely map to multiple locations with equal scores and are reported as multi-mapping alignments. This is expected and correct behaviour: the multi-mapping rate is a property of the genome annotation, not an aligner artefact.

8.6 Adjunct-introns: splice-boundary optimisation

After the dynamic-programming pass, each exon fragment still carries its raw jumper-extension coordinates and has not been clipped at splice boundaries. The adjunct-introns step processes every successive pair of exon fragments along a path.

Intron size threshold. Magic2 does not call introns below 30 bases. No confirmed human intron shorter than 30 bases is known, and short genomic gaps in an alignment are more parsimoniously explained by sequencing deletions: a 6-base gap is reported as two successive 3-base deletions and penalised accordingly in the score. Above 30 bases, a gap becomes an intron candidate, but confidence is graded rather than binary. Magic2 counts, for each candidate gap-length threshold t , the fraction of gaps at or above t that carry a canonical GT-AG (or CT-AC) signal. As t increases this fraction rises; the threshold is set at the point where the slope of that curve falls to 1 (i.e. where raising the threshold further gains no additional enrichment for canonical signals). This adaptive procedure is described in detail in the Magic-BLAST paper (?); Magic2 inherits and extends it. An intron is considered reliable if it either carries a GT-AG signal or is supported by multiple independent reads.

Intron scoring bonus/penalty. Once a gap is confirmed as an intron, its contribution to the alignment score is adjusted by one additional `errCost` unit depending on the library type inferred during the pilot phase (Section 4). In RNA mode the intron receives a bonus of `+errCost`: a spliced alignment is rewarded relative to an unspliced one, favouring genuine gene models. In DNA mode the intron receives a penalty of `-errCost`: a genomic gap is treated with mild suspicion, which disfavours retroposon insertions and other spurious long-range matches that would otherwise masquerade as spliced alignments. The net effect is that the same scoring framework handles both RNA-seq and DNA-seq libraries correctly without any mode-specific code paths. Two adjacent exon fragments often overlap in read coordinates: both extensions reached into the intron and claimed the same read bases. The overlap is resolved by choosing a clip point — a position in read space at which the left fragment ends and the right fragment begins — that minimises the total number of errors (substitutions plus indels of up to 3 bases) summed across both fragments. In low-complexity regions the overlap can span 30 or more bases with errors on both sides, so this joint minimisation is essential for accurate boundary placement.

Splice-signal selection. Once the optimal clip point is identified, the genomic gap between the two fragments defines a candidate intron. If a junction seed hit is available for this locus, its annotated boundary is adopted directly without further search. Otherwise, Magic2 slides the clip point within a small window around the error minimum and scores the dinucleotides at the resulting donor and acceptor sites, applying the following preference order:

1. GT-AG (canonical sense) or CT-AC (canonical antisense), chosen according to the strand established during the pilot phase (Section 4);
2. GC-AG (less common but well-documented);
3. any other dinucleotide pair, scored by error count alone.

For non-stranded libraries both GT-AG and CT-AC are considered equally at step 1, and the one consistent with the lower error count is preferred. The final clip point therefore simultaneously minimises alignment errors and maximises splice-signal evidence.

8.7 Adjunct-exons: shadow realignment

After intron boundaries are fixed, the path coordinates are used to extract the *genomic shadow* of the read: the exact genomic sequence spanned by the alignment, with introns removed. The full read is then realigned against this personalised reference using the jumper algorithm.

This second pass resolves two classes of residual error. First, a short gap in the read that was previously interpreted as a sequencing indel may, when viewed against the shadow, be recognisable as a mis-measured intron boundary; adding two to five missing genomic bases from the better-supported side can create a clean GT-AG signal that was invisible in the original alignment. Second, errors near exon boundaries are sometimes artefacts of the clip point rather than true mismatches; realignment against the shadow redistributes them correctly.

After shadow realignment, all mismatches and indels are re-projected into exonic read coordinates, and the final alignment score is computed (Section 8.8). This is the score used for all downstream comparisons.

8.8 Alignment scoring

The score of a single-read alignment path is:

$$s = A - \text{errCost} \times E$$

where A is the number of aligned bases in read coordinates (each read base counted once, regardless of how many exons it spans), E is the total number of errors (substitutions plus indels of 1–3 bases), and $\text{errCost}=4$ is the default penalty per error. Genomic gaps of 4–29 bases are treated as deletions and contribute to E (a 6-base gap counts as two successive 3-base deletions, penalised as $2 \times \text{errCost}$). Gaps of 30 bases or more are candidates for introns and do not contribute to E ; their status is confirmed by splice-signal evidence and support count as described in Section 8.6.

8.9 Pair scoring and multi-mapping resolution

For paired-end libraries, the two mates of a pair are scored jointly. If both mates align at the same locus with correct relative orientation and a plausible insert size, their individual scores are summed to form a *pair score*. The pair score is then compared across all candidate loci for that read pair.

Consider a pair that aligns at two loci (both mates present at each) and also has an orphan alignment of one mate at a third locus. The orphan is discarded: a validated pair score always dominates a single-mate score, because the joint alignment provides roughly twice the evidence. All loci whose pair score equals the best pair score are retained as equally supported placements; this set is reported as the multi-mapping output for that pair.

For single-end libraries there are no pair scores; the best individual score (or all tied-best scores) is retained.

Reads or pairs whose best score falls below $\text{minScore}=30$ or whose aligned length is below $\text{minAli}=30$ are discarded.

8.10 Long-read rafias

Nanopore and PacBio reads are long enough to span many exons in a single read, but their error rate is substantially higher than Illumina and their errors are not uniformly distributed: certain sequence motifs,

homopolymer runs, and locally low-complexity regions can produce a stretch of 30 to several hundred bases that cannot be reliably matched to the genome. In such cases the dynamic-programming step may produce two separate paths for what is biologically a single continuous alignment: for example, exons 1–7 on the left and exons 9–14 on the right, with a gap in between that no seed covers.

If the two paths share a unique genomic locus, have consistent orientation, and the unmatched read region is shorter than the genomic gap between the two flanking exon groups, Magic2 merges them into a single structure called a *rafia*. The unmatched central portion is called the *bubble*. The bubble is assumed to correspond to a real genomic region between the flanking exons — perhaps a difficult-to-sequence segment or a region absent from the current genome assembly — but its precise internal coordinates are unknown.

Scoring. Bubble bases are not penalised as substitutions. The alignment score is computed from the matched flanking exons only, exactly as for a standard path (Section 8.8). This prevents a long unsequenceable stretch from artificially depressing the score of an otherwise high-quality alignment.

BAM output. No single standard CIGAR operator can represent a bubble: the read bases are present and real, but their precise genomic location is unknown within a large reference span. Magic2 encodes a rafia using a compound `I/N` pair placed at the bubble boundary:

```
700M200I3000N500M
```

where `700M` and `500M` stand for the left and right exon groups (each expanding to their own `M/I/D` operations in practice), `200I` records the 200 bubble bases as present in the read but not consumed from the reference, and `3000N` records the 3000-base reference span that is skipped. The adjacency of `200I3000N` signals that these two events are coupled: the 200 read bases are not a sequencing insertion but an unlocatable fragment sitting somewhere within the 3000-base genomic region. The full read sequence is preserved in the `SEQ` field across all bases, so scoring and downstream re-analysis are unaffected.

This encoding is syntactically valid under the SAM specification, which does not forbid adjacent `I` and `N` operators. Tools that compute read length as the sum of read-consuming operators will recover the correct length. Some tools that interpret `I` strictly as a short sequencing insertion may flag the record; this is a known limitation of the BAM format for long-read structural variants.

[TODO — implement and verify this CIGAR encoding for rafias in `wsa/sa.sam.c`, and test against samtools, IGV, and any downstream tools used in the Magic2 pipeline.]

Rafias are rarely encountered in Illumina data: short reads that hit an unreadable stretch typically fail to align altogether rather than producing two flanking paths.

9 Adapter detection and trimming

Adapter sequences may be supplied via `--adaptor1/--adaptor2` or the run manifest, but this is not required. In practice, the exact adapter sequence as it appears in sequenced reads is frequently unknown to the user, owing to the complexity of modern library constructions. Magic2 therefore infers adapter sequences directly from the data.

9.1 Overhang accumulation

After alignment, the unaligned 3' tail of a read — the overhang — is a candidate adapter fragment. Magic2 accumulates four per-position base-frequency histograms over overhangs, one for each possible first overhang base (A, T, G, C). The split is necessary because the first overhang base has a 25 % probability of matching

the genomic sequence by chance, which otherwise shifts the composite histogram and obscures the adapter signal.

To reconstruct the adapter consensus, the histogram with the highest total count is selected (say, the A histogram, implying the adapter begins with A). The most frequent second base is then identified (say, G), and the G sub-histogram, shifted by one position, is added to the A histogram. The resulting composite histogram is extremely clean: on real datasets, base prevalence at each of the first ~40 adapter positions typically exceeds 70 % (Figure 1).




Figure 1: Per-position base-frequency of 3' unaligned overhangs (run A_Total_151PE), after splitting by first overhang base and compositing. **(a)** Read 1 (TruSeq adapter): the four bases alternate sharply between positions, each dominant base reaching 88–94 %, unambiguously encoding the 29-nt consensus shown above the panel. **(b)** Read 2 (degenerate adapter): dominant-base fractions range from 50–80 %, consistent with a partially randomised ligation sequence; the adapter family is nonetheless uniquely identified. Colours: A = green, T = red, G = orange, C = blue. The coloured sequence above each panel is the adapter inferred by Magic2.

9.2 Pilot blocks and reprocessing

Adapter inference is part of the broader pilot-block mechanism described in Section 4. Once the adapter consensus is called, any pilot block that was processed without it is reprocessed from scratch; no reads are lost.

9.3 Trimming and alignment extension

Once an adapter is known, each read's overhang is compared to it. If enough bases match, the alignment is clipped back to the exact genomic position where the adapter begins, effectively extending the reported alignment. In the Magic2 output format the trimmed bases are exported as clipped, and the extension bases are reported in upper case rather than lower case. A read that aligns fully up to the recognised adapter boundary, with no mismatches in the genomic portion, is counted as a *perfect alignment*.

9.4 RNA-mode extras: poly(A) detection

RNA mode and strand bias are inferred during the pilot phase (Section 4). Once the strand is established, a poly(A) tail is sought at the 3' end of forward reads and a poly(T) tail at the 5' end of reverse reads, the two being equivalent by symmetry. Tails carrying at least eight consecutive A (or T) residues are recorded, accumulated across reads, and reported as candidate poly(A) addition sites.

9.5 Trans-splice leader detection (*C. elegans*)

In *C. elegans* the majority of mRNAs acquire a short, stereotyped sequence at their 5' end by trans-splicing: the *splice leader* (SL). SL sequences are highly conserved and serve as extremely reliable markers of transcription start sites. Magic2 searches for SL sequences at the same positions used for adapter detection: the 5' end of forward reads and the 3' end of reverse reads. The 12 known *C. elegans* SL sequences (SL1–SL12) are hard-coded in the program (?). Reads carrying a recognised SL are trimmed at the SL boundary, the SL identity is recorded, and the counts are reported in the TSF output, providing a direct census of trans-splicing activity across the transcriptome.

SL detection is activated by specifying the target organism in the run configuration. The recommended mechanism is a `Species=worm` entry in the `-T` target configuration file, which keeps species-specific settings alongside the genome index rather than on the command line. A `--species worm` command-line override is also accepted for convenience.

10 Intron boundary optimisation and splice-signal scoring

Intron detection, boundary optimisation, splice-signal scoring, the dynamic intron-size threshold, and the RNA/DNA scoring adjustment are all handled during the adjunct-introns step described in Section 8.6. The maximum reportable intron size is `maxIntron=1,000,000` bases; gaps above this limit are not called as introns regardless of splice-signal evidence. Strand inference from the GT-AG versus CT-AC intron balance is performed during the pilot phase (Section 4) and used throughout.

11 Output formats

11.1 Alignment: SAM and BAM

Magic2 writes alignments in SAM or BAM format (`wsa/sa.sam.c`). The SAM writer is self-contained: it produces a fully valid SAM file with no dependency on any external library. BAM output is requested via

`--bam`: Magic2 checks whether `samtools view` is available in `PATH` and, if so, pipes its SAM output through `samtools view -bS` to produce a compressed BAM file. If `samtools` is not found, Magic2 falls back silently to uncompressed SAM and reports the fallback to the user. This keeps the binary free of any link-time dependency on `htslib` or `samtools`.

11.2 Compact hits format

The `.hits` format (`wsa/sa.tabular.c`), selected via `--hits`, is a compact tab-separated alignment format designed for internal Magic2 pipelines where BAM compatibility is not required. It is faster to write and parse than SAM and represents multi-exon alignments more compactly.

11.3 TSF: Tag-Sample Format

All scalar outputs — expression counts, intron support tables, poly(A) sites, QC metrics, adaptor profiles — are exported in TSF (Tag-Sample Format): a plain tab-separated file with columns `run / tag / format-code / value`, where format codes are `i` (integer), `f` (float), and `t` (text). TSF files need not be sorted; multiple TSF files can be concatenated directly to merge runs; and the `tsf.c` library can pivot any TSF file into a two-dimensional table compatible with Excel.

11.4 Coverage tracks (wiggle / BigWig)

Coverage profiles are exported strand-specifically by `wsa/sa.wiggle.c` in the UCSC *fixedStep* wiggle format (`.BF.gz`), which is the most compact text wiggle encoding. Each chromosome block opens with a header line of the form

```
fixedStep chrom=chr1 start=1 step=10
```

followed by one coverage value per line per step position. The step size is chosen automatically from genome length: 1 for small genomes (bacteria, yeast) and 10 for large genomes (human, mouse), consistent with Magic2's general policy of deriving all internal parameters from the data rather than requiring user input. The step can be overridden with `--wiggle_step N` when a specific resolution is needed. Read-start and read-end profiles are also available for transcription-start-site (TSS) and transcription-end-site (TES) analysis.

The current standard for genome browsers (UCSC, IGV, Ensembl) is BigWig: an indexed binary format. BigWig cannot be produced natively without linking the UCSC Kent library, which Magic2 deliberately avoids. The recommended workflow is therefore identical to the BAM strategy: Magic2 checks whether `wigToBigWig` is available in `PATH` and, if so, pipes the *fixedStep* wiggle output through it, also writing the required chromosome-sizes file from the genome index. If the tool is absent, the compressed *fixedStep* file is kept. `wigToBigWig` is distributed as a standalone static binary by UCSC (<https://hgdownload.cse.ucsc.edu/admin/exe/>) and is also installable via `conda install ucsc-wigtobigwig`.

[TODO: implement the `wigToBigWig` pipe and chromosome-sizes export in `wsa/sa.wiggle.c`.]

11.5 Intron support tables

Intron evidence is exported in two TSF files per run.

Single-intron table (`introns.tsf`). Each record identifies one intron and one run. The tag field encodes the intron as `chrom_donor_acceptor`, where the double underscore separates the chromosome name from the two boundary coordinates and a single underscore separates donor from acceptor. Negative coordinates indicate the antisense strand. The format code is `iit` (two integers, one text): `n` (number of supporting reads on the annotated strand), `nR` (reads on the opposite strand), and `feet` (splice signal, e.g. `gt_ag`, `ct_ac`, `gc_ag`). A typical record is:

```
chr1__146546_138479 A_Total_151PE iit 4 0 gt_ag
```

The chromosome name is taken verbatim from the genome FASTA file and the optional GTF annotation; it appears unchanged in all output files, so coordinates are unambiguous across tools.

Double-intron table (`double_introns.tsf`). Each record reports a pair of introns observed within the same read. This is more informative than single-intron support alone: a read spanning two consecutive introns confirms the intervening exon as a genuine transcribed unit, providing direct exon-connectivity evidence for transcript reconstruction. The tag encodes both introns as `chrom__a1_a2__b1_b2` (triple underscore between the two intron coordinate pairs). The format code is `iitt: n, nR, feet1, feet2`.

Both files are directly convertible to Excel-compatible tables via:

```
tsf -i introns.tsf -I tsf -O table -o introns_table.txt
```

and can be merged across runs by simple concatenation, since TSF files require no sorting.

11.6 Poly(A) addition sites

Candidate poly(A) addition sites are exported in `polyA.tsf`. Each record identifies one site, one run, and one strand. The tag field has the form `G.chrom__position`, where `G` denotes a genomic locus (as opposed to a gene model coordinate), the chromosome name is verbatim from the FASTA, and the position is the last aligned base before the poly(A) tail. The format code is `ti` (one text, one integer): `Strand (+ or -)` and `n` (number of reads with a poly(A) or poly(T) tail ending at this position). A typical record is:

```
G.chr1__1304288 A_Total_151PE ti + 5
```

Sites supported by only one read are retained in the file but should be treated with caution; reliable poly(A) sites typically accumulate tens to hundreds of reads at a single nucleotide position, reflecting the precision of cleavage-and-polyadenylation.

11.7 Error profile

The per-run substitution and indel profile is exported in `errorProfile.tsf`. The format code is `i` throughout; the tag encodes the error type directly. The 12 substitution types are named `X>Y` (reference base to observed base, e.g. `A>G`); diagonal entries (`A>A` etc.) are always zero by construction and serve as a consistency check. Insertions and deletions of one base are counted separately by the inserted or deleted base (`InsA`, `DelT`, etc.). `Any` gives the total error count across all categories. A typical file:

```
A_Total_151PE Any i 4880463
A_Total_151PE A>G i 299012
A_Total_151PE G>T i 581599
A_Total_151PE InsA i 73057
A_Total_151PE DelA i 52330
```

This profile characterises the sequencing technology rather than the biology: the balance between transitions (`A>G`, `G>A`, `C>T`, `T>C`) and transversions (`A>C`, `A>T`, `G>C`, `G>T`, `C>A`, `C>G`, `T>A`, `T>G`) in the mismatch table reflects the error chemistry of the sequencer, not the underlying SNP or mutation spectrum. The profile is useful for comparing sequencing quality across runs, for detecting systematic chemistry artefacts, and for tuning the `errCost` parameter for non-Illumina technologies.

11.8 Run statistics

The file `runStats.tsf` collects all global alignment statistics for a run in a single TSF file (`wsa/sa.stats.c`). The header comment records the index used, seed length, maximum repeats, and error cost, providing a self-contained audit trail.

Unlike the other TSF files, `runStats.tsf` uses multi-value format codes (`ift`, `iftift`, `itititft`, ...) that pack several related numbers and labels into one record. This keeps the file human-readable as a flat list while allowing the `tsf` tool to expand it into a table. Tags with zero counts are sometimes omitted.

The main categories are as follows.

Pair and read counts. `Pairs`, `Aligned_pairs`, `Compatible_pairs` (correct orientation and insert size), `Circle_pairs` (both mates on the same strand — a library artefact), and `Non_compatible_pairs`. At read level: `Aligned_reads`, `Unaligned_reads`, `Perfect_reads` (zero mismatches, zero indels), `Partial_reads` (one mate aligned), and `Complex_reads` (unusual alignment structure). `Low_entropy_reads` counts reads filtered out before alignment due to insufficient sequence complexity.

Base counts. `Bases` gives total, R1, and R2 base counts as three integers (`iii`). `Aligned_bases` gives total, R1, and R2 aligned counts each with their percentage of raw bases (`ifift`).

Error summary. `Errors` gives the total error count and its percentage of aligned bases. `Mismatches_and_InDels` gives the subset of errors that are substitutions or short indels (1–3 bases), excluding intron-spanning gaps. The full per-type breakdown is in `errorProfile.tsf` (Section 11).

Adapter and clipping statistics. `Clipped_3prime_Adaptor_read1` records the number of reads trimmed and the inferred adapter sequence. `Clipped_polyA` and `Clipped_polyT` count reads trimmed at a poly(A) or poly(T) tail.

Intron statistics. `Intron_supports` gives the total number of intron-spanning reads, broken down by splice signal (GT-AG, CT-AC, other) and the fraction on the positive strand. `Supported_introns` gives the number of distinct intron loci passing the reliability threshold (GT-AG seen once, or any signal seen at least three times).

Multi-mapping distribution. `Reads_Aligned_once` through `Reads_multi_aligned_10` give the number and percentage of reads aligning to exactly 1, 2, ..., 10 loci. `Reads_aligned_several_times` is the total fraction with more than one alignment.

Target-class breakdown. `Magic2` aligns against one or more named target classes defined in the index (Section 5). The standard classes are G (genomic, the primary genome), M (mitochondrial), and R (repeat / ribosomal RNA). For each class, `Reads_aligned_in_class_X` and `Bases_aligned_in_class_X` give total counts, positive- and negative-strand breakdowns, the positive-strand fraction, and the percentage of all aligned reads or bases. The strand balance within each class is informative: in a stranded RNA-seq library, class G reads show roughly 50/50 strand split across the whole genome but strong strand bias within individual genes, while class R reads are nearly all on one strand because ribosomal RNA transcription is unidirectional.

12 Automatic parameter selection

Magic2 derives all internal parameters from the hardware and data rather than requiring user configuration. Hardware auto-configuration (NUMA node selection, CPU count, agent and block counts) is described in Section 3. Data auto-configuration (read type, library type, strand, adapter sequences) is described in Section 4. The resulting defaults are listed below for reference; all can be overridden on the command line when needed.

Parameter	Default
errCost	4
errRateMax	12 %
bMax	3 Mb per block
Index step	6 (genomes > 1 Mb)
Alignment step (iStep)	tStep/2
seedLength	16 (genomes < 1 Gb); 18 (human/mouse)
maxTargetRepeats	81 (Section 5.5)
wiggle_step	1 (small genomes); 10 (human/mouse)
minAli	30
minScore	30
maxIntron	1,000,000 bases

13 Build system and supported platforms

Magic2 is written in C and built with `make` and `tcsh`. The repository lives inside the global `acedb` distribution (~700,000 lines of C); Magic2 depends on `acedb` libraries (notably `w1`, `wego`) and has not yet been fully separated.

Setting	Value
Optimised Linux build flag	ACEDB_MACHINE=LINUX_4_OPT
64-bit variant	ACEDB_MACHINE=LINUX_X86_64_4_OPT
Output binary	Magic2/bin.LINUX_4_OPT/magic2
Primary platform	Linux x86-64
Secondary platforms	Linux ARM, macOS
Optional GPU acceleration	CUDA (Thrust); silent CPU fallback

Continuous integration currently uses Travis CI (`.travis.yml`). Migration to GitHub Actions with matrix builds for all three target platforms is planned.

14 Current status and roadmap

14.1 Implemented and tested

The following components are fully implemented and have been validated on human, mouse, and *C. elegans* RNA-seq datasets: sort-merge alignment, radix sort, Go-style channel parallelism, NUMA-aware respawning, pilot-block inference, adapter detection and trimming, splice-boundary optimisation (adjunct-introns, Section 8.6), shadow realignment (adjunct-exons, Section 8.7), poly(A) detection, trans-splice leader detection, SAM output, TSF output, fixedStep wiggle output, intron and double-intron tables, error profile, and run statistics. Optional GPU acceleration via CUDA is implemented and tested on NVIDIA hardware; execution falls back silently to the CPU path on machines without a compatible device.

14.2 Implemented, verification pending

Two output features are coded but require testing against standard downstream tools before the first public release:

- **Rafia BAM encoding.** The compound 200I3000N CIGAR strategy for long-read bubbles (Section 8.10) needs to be implemented and verified against samtools, IGV, and downstream tools (see `wsa/sa.sam.c`).
- **BigWig export.** The `wigToBigWig` pipe and automatic chromosome-sizes file (Section 11.4) need to be implemented in `wsa/sa.wiggle.c`.

14.3 Planned

The following are design decisions made but not yet implemented:

- GitHub Actions CI with matrix builds for Linux x86-64, Linux ARM, and macOS, replacing the current Travis CI configuration.
- Automated GitHub Releases attaching pre-built binaries and documentation when a version tag is pushed.
- A regression test suite running a small reference dataset and comparing outputs to stored expected results.
- Verification that Magic2 runs correctly inside Docker containers (currently untestable at NCBI for security reasons).
- A standalone TSF library reference (`magic2_tsf_reference.tex`).

References

- Boratyn, G. M., Thierry-Mieg, J., Thierry-Mieg, D., Busby, B. & Madden, T. L. Magic-BLAST, an accurate RNA-seq aligner for long and short reads. *BMC Bioinformatics* **20**, 405 (2019).
- Thierry-Mieg, D. & Thierry-Mieg, J. AceView: a comprehensive cDNA-supported gene and transcripts annotation. *Genome Biology* **7**, S12 (2006).
- Buchfink, B., Reuter, K. & Drost, H.-G. Sensitive protein alignments at tree-of-life scale using DIAMOND. *Nature Methods* **18**, 366–368 (2021).
- Knuth, D. E. *The Art of Computer Programming, Vol. 3: Sorting and Searching* (Addison-Wesley, 1973).
- Frigo, M., Leiserson, C. E., Prokop, H. & Ramachandran, S. Cache-oblivious algorithms. In *Proceedings of the 40th Annual Symposium on Foundations of Computer Science*, 285–297 (1999).
- Donovan, A. A. A. & Kernighan, B. W. *The Go Programming Language* (Addison-Wesley Professional, 2015).
- Hoare, C. A. R. Communicating sequential processes. *Communications of the ACM* **21**, 666–677 (1978).
- Blumenthal, T. *et al.* A global analysis of *Caenorhabditis elegans* operons. *Nature* **417**, 851–854 (2002).